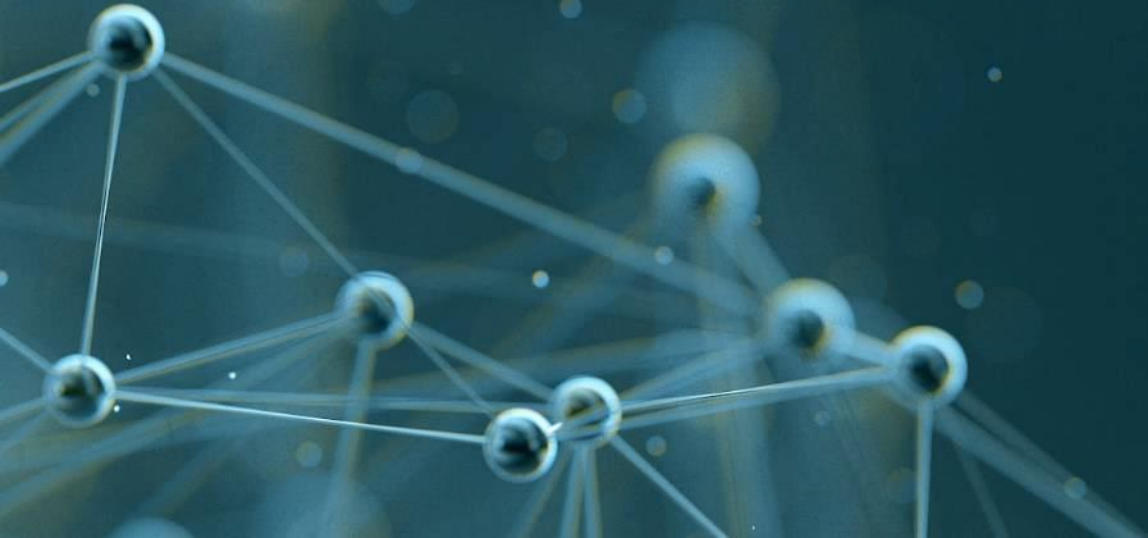


# BOOSTING AND STACKING



# REVIEW OF BAGGING

Grow decision tree from multiple bootstrapped samples.

Date	Title	Budget	DomesticTotalGross	Director	Rating	Runtime
0	2013-11-22 The Hunger Games: Catching Fire	130000000	424688047	Francis Lawrence	PG-13	146
1	2013-05-03 Iron Man 3	200000000	409013994	Shane Black	PG-13	129
2	2013-11-22 Frozen	150000000	400738009	Chris BuckJennifer Lee	PG	108
3	2013-07-03 Despicable Me 2	76000000	368861265	Pierre CoffinChris Renaud	PG	98
4	2013-06-14 Man of Steel	225000000	291045518	Zack Snyder	PG-13	143
5	2013-10-04 Gravity	100000000	274082705	Alfonso Cuaron	PG-13	91
6	2013-06-21 Monsters University	NaH	268482764	Dan Scanlon	G	107
7	2013-12-13 The Hobbit: The Desolation of Smaug	NaH	258366555	Peter Jackson	PG-13	161
8	2013-05-24 Fast & Furious 6	160000000	238679850	Justin Lin	PG-13	130
9	2013-03-08 Oz the Great and Powerful	215000000	234911825	Sam Raimi	PG	127
10	2013-05-16 Star Trek Into Darkness	190000000	228718661	J.J. Abrams	PG-13	123
11	2013-11-08 Thor: The Dark World	170000000	206362140	Alan Taylor	PG-13	120
12	2013-06-21 World War Z	190000000	20258711	Marc Forster	PG-13	116
13	2013-03-22 The Croods	135000000	187168425	Kirk De MiccoChris Sanders	PG	98
14	2013-06-28 The Heat	43000000	159582188	Paul Feig	R	117
15	2013-08-07 We're the Millers	37000000	150384119	Rawson Marshall Thurber	R	110
16	2013-12-13 American Hustle	40000000	150117807	David O. Russell	R	138
17	2013-05-10 The Great Gatsby	105000000	144840419	Baz Luhrmann	PG-13	143

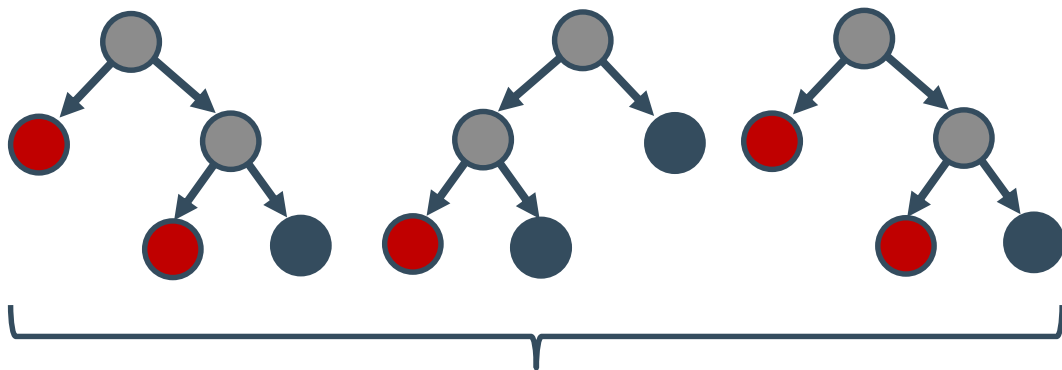
Date	Title	Budget	DomesticTotalGross	Director	Rating	Runtime
0	2013-11-22 The Hunger Games: Catching Fire	130000000	424688047	Francis Lawrence	PG-13	146
1	2013-05-03 Iron Man 3	200000000	409013994	Shane Black	PG-13	129
2	2013-11-22 Frozen	150000000	400738009	Chris BuckJennifer Lee	PG	108
3	2013-07-03 Despicable Me 2	76000000	368861265	Pierre CoffinChris Renaud	PG	98
4	2013-06-14 Man of Steel	225000000	291045518	Zack Snyder	PG-13	143
5	2013-10-04 Gravity	100000000	274082705	Alfonso Cuaron	PG-13	91
6	2013-06-21 Monsters University	NaH	268482764	Dan Scanlon	G	107
7	2013-12-13 The Hobbit: The Desolation of Smaug	NaH	258366555	Peter Jackson	PG-13	161
8	2013-05-24 Fast & Furious 6	160000000	238679850	Justin Lin	PG-13	130
9	2013-03-08 Oz the Great and Powerful	215000000	234911825	Sam Raimi	PG	127
10	2013-05-16 Star Trek Into Darkness	190000000	228718661	J.J. Abrams	PG-13	123
11	2013-11-08 Thor: The Dark World	170000000	206362140	Alan Taylor	PG-13	120
12	2013-06-21 World War Z	190000000	20258711	Marc Forster	PG-13	116
13	2013-03-22 The Croods	135000000	187168425	Kirk De MiccoChris Sanders	PG	98
14	2013-06-28 The Heat	43000000	159582188	Paul Feig	R	117
15	2013-08-07 We're the Millers	37000000	150384119	Rawson Marshall Thurber	R	110
16	2013-12-13 American Hustle	40000000	150117807	David O. Russell	R	138
17	2013-05-10 The Great Gatsby	105000000	144840419	Baz Luhrmann	PG-13	143

Date	Title	Budget	DomesticTotalGross	Director	Rating	Runtime
0	2013-11-22 The Hunger Games: Catching Fire	130000000	424688047	Francis Lawrence	PG-13	146
1	2013-05-03 Iron Man 3	200000000	409013994	Shane Black	PG-13	129
2	2013-11-22 Frozen	150000000	400738009	Chris BuckJennifer Lee	PG	108
3	2013-07-03 Despicable Me 2	76000000	368861265	Pierre CoffinChris Renaud	PG	98
4	2013-06-14 Man of Steel	225000000	291045518	Zack Snyder	PG-13	143
5	2013-10-04 Gravity	100000000	274082705	Alfonso Cuaron	PG-13	91
6	2013-06-21 Monsters University	NaH	268482764	Dan Scanlon	G	107
7	2013-12-13 The Hobbit: The Desolation of Smaug	NaH	258366555	Peter Jackson	PG-13	161
8	2013-05-24 Fast & Furious 6	160000000	238679850	Justin Lin	PG-13	130
9	2013-03-08 Oz the Great and Powerful	215000000	234911825	Sam Raimi	PG	127
10	2013-05-16 Star Trek Into Darkness	190000000	228718661	J.J. Abrams	PG-13	123
11	2013-11-08 Thor: The Dark World	170000000	206362140	Alan Taylor	PG-13	120
12	2013-06-21 World War Z	190000000	20258711	Marc Forster	PG-13	116
13	2013-03-22 The Croods	135000000	187168425	Kirk De MiccoChris Sanders	PG	98
14	2013-06-28 The Heat	43000000	159582188	Paul Feig	R	117
15	2013-08-07 We're the Millers	37000000	150384119	Rawson Marshall Thurber	R	110
16	2013-12-13 American Hustle	40000000	150117807	David O. Russell	R	138
17	2013-05-10 The Great Gatsby	105000000	144840419	Baz Luhrmann	PG-13	143

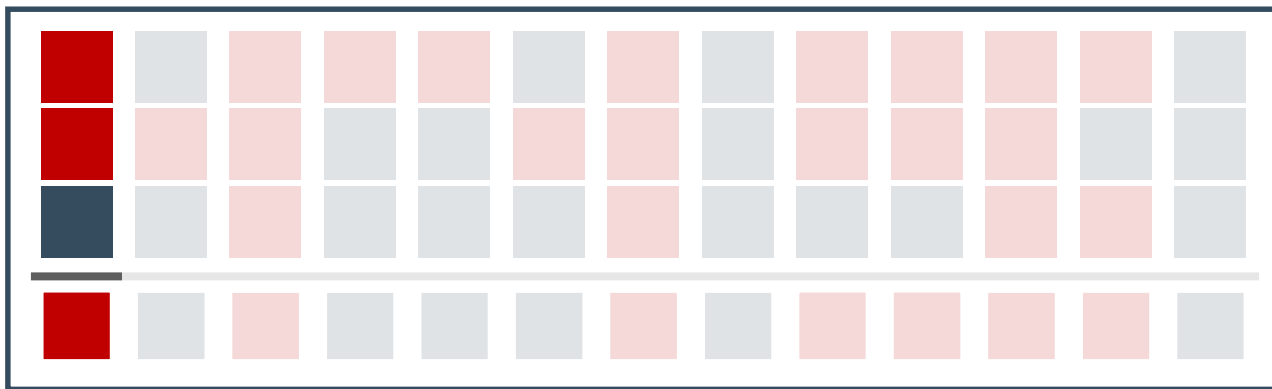
Date	Title	Budget	DomesticTotalGross	Director	Rating	Runtime
0	2013-11-22 The Hunger Games: Catching Fire	130000000	424688047	Francis Lawrence	PG-13	146
1	2013-05-03 Iron Man 3	200000000	409013994	Shane Black	PG-13	129
2	2013-11-22 Frozen	150000000	400738009	Chris BuckJennifer Lee	PG	108
3	2013-07-03 Despicable Me 2	76000000	368861265	Pierre CoffinChris Renaud	PG	98
4	2013-06-14 Man of Steel	225000000	291045518	Zack Snyder	PG-13	143
5	2013-10-04 Gravity	100000000	274082705	Alfonso Cuaron	PG-13	91
6	2013-06-21 Monsters University	NaH	268482764	Dan Scanlon	G	107
7	2013-12-13 The Hobbit: The Desolation of Smaug	NaH	258366555	Peter Jackson	PG-13	161
8	2013-05-24 Fast & Furious 6	160000000	238679850	Justin Lin	PG-13	130
9	2013-03-08 Oz the Great and Powerful	215000000	234911825	Sam Raimi	PG	127
10	2013-05-16 Star Trek Into Darkness	190000000	228718661	J.J. Abrams	PG-13	123
11	2013-11-08 Thor: The Dark World	170000000	206362140	Alan Taylor	PG-13	120
12	2013-06-21 World War Z	190000000	20258711	Marc Forster	PG-13	116
13	2013-03-22 The Croods	135000000	187168425	Kirk De MiccoChris Sanders	PG	98
14	2013-06-28 The Heat	43000000	159582188	Paul Feig	R	117
15	2013-08-07 We're the Millers	37000000	150384119	Rawson Marshall Thurber	R	110
16	2013-12-13 American Hustle	40000000	150117807	David O. Russell	R	138
17	2013-05-10 The Great Gatsby	105000000	144840419	Baz Luhrmann	PG-13	143

# REVIEW OF BAGGING

Vote on or average  
result from each tree  
for each data point.

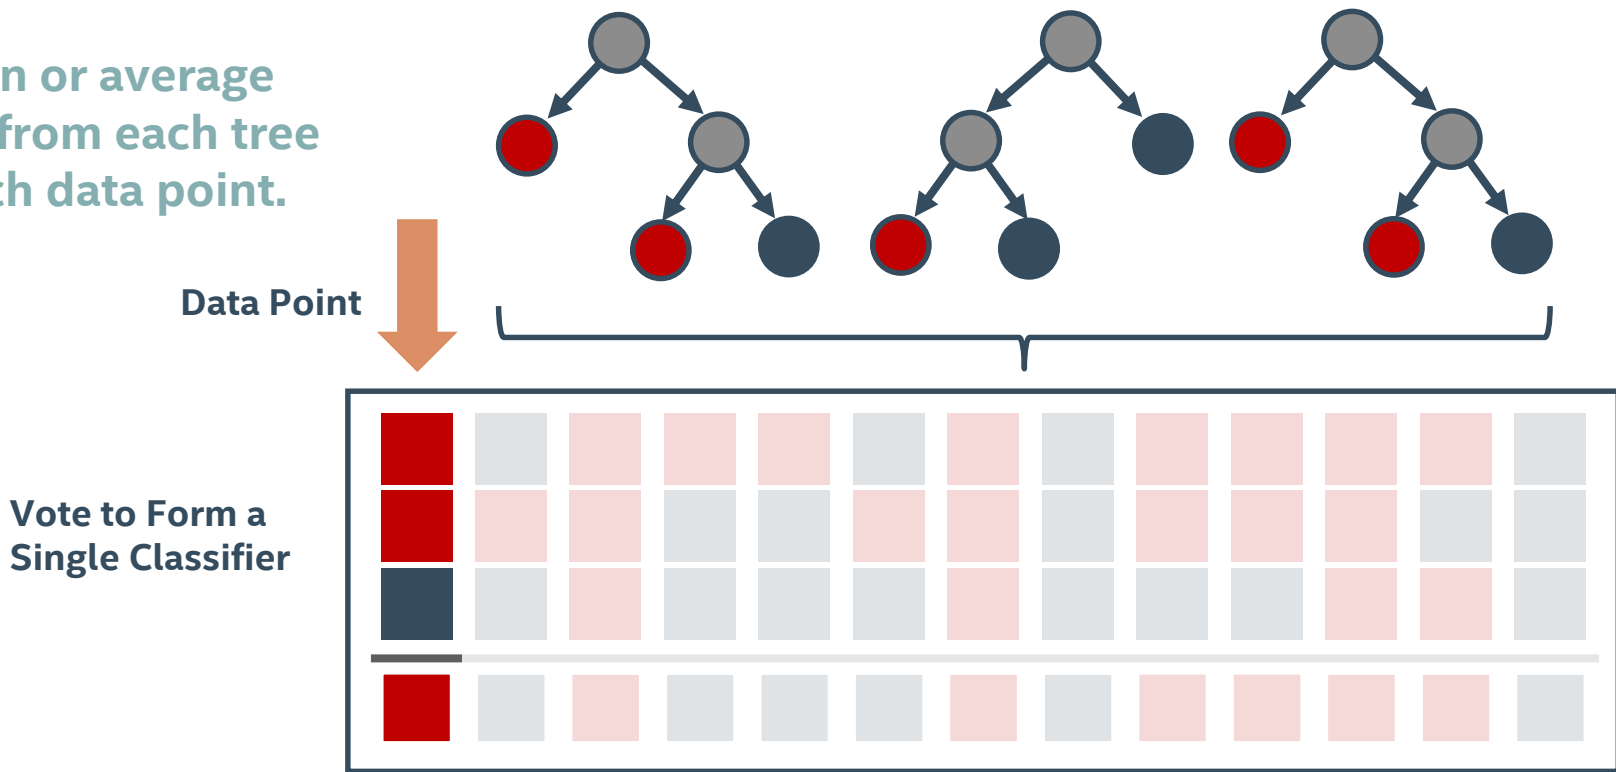


Vote to Form a  
Single Classifier



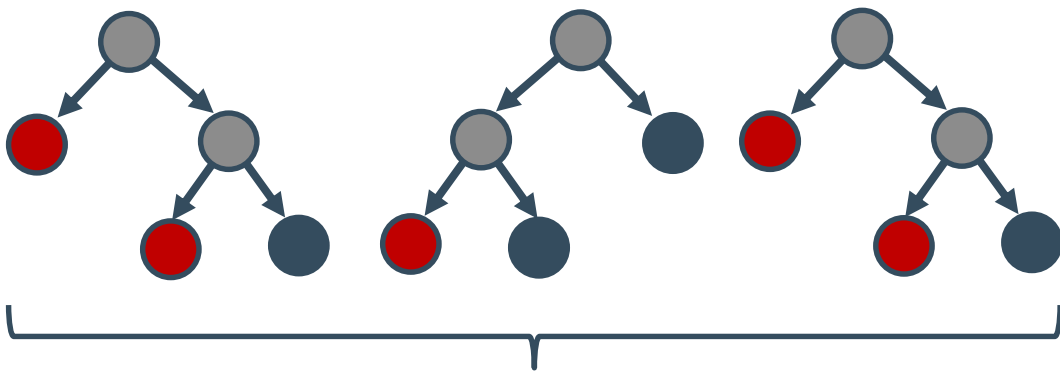
# REVIEW OF BAGGING

Vote on or average  
result from each tree  
for each data point.



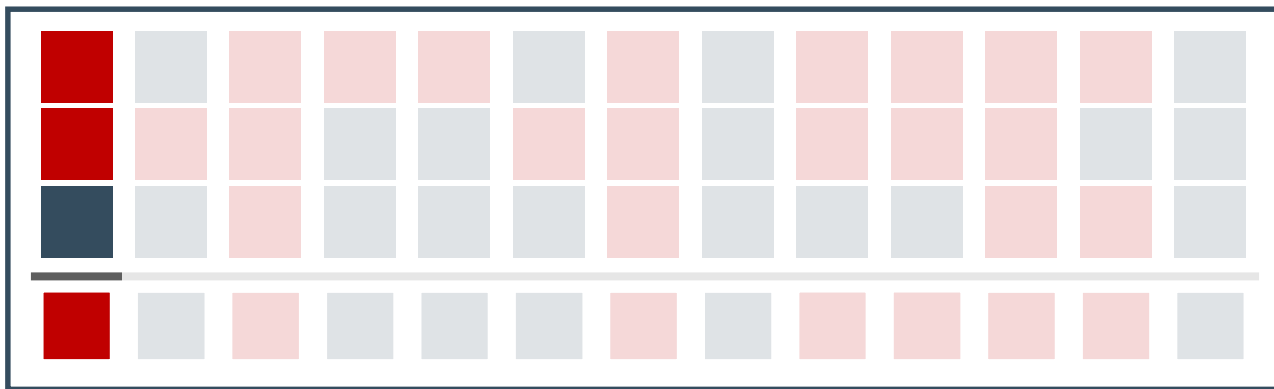
# REVIEW OF BAGGING

Vote on or average  
result from each tree  
for each data point.



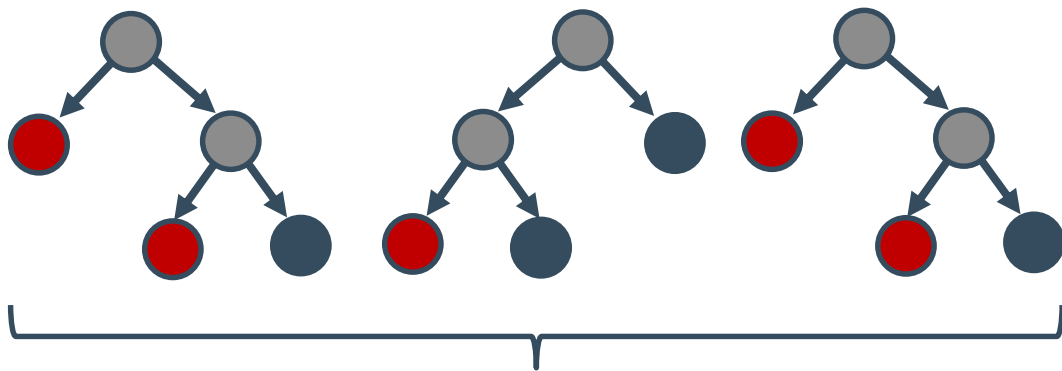
Vote to Form a  
Single Classifier

Results

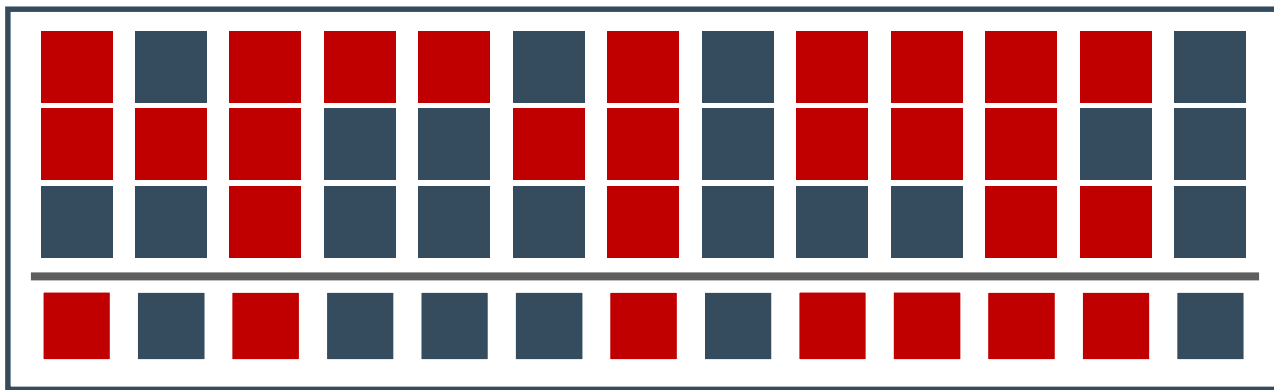


# REVIEW OF BAGGING

Vote on or average  
result from each tree  
for each data point.

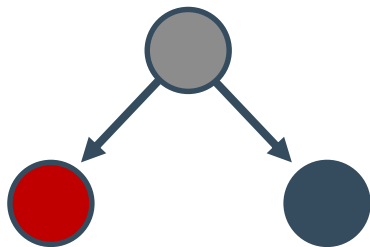


Vote to Form a  
Single Classifier



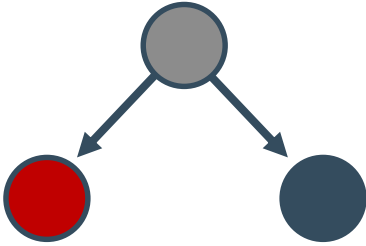
# DECISION STUMP: THE BOOSTING BASE LEARNER

Temperature  $> 50^{\circ}\text{F}$

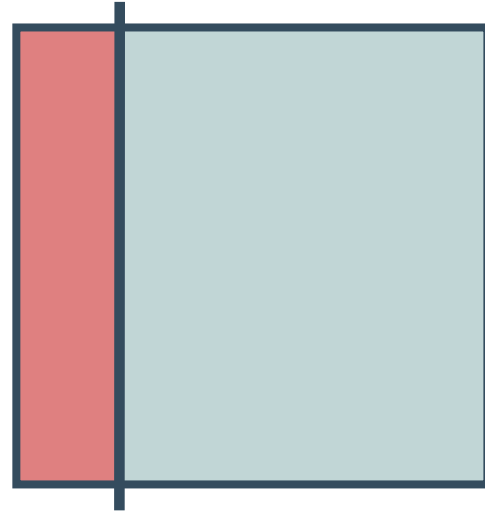


# DECISION STUMP: THE BOOSTING BASE LEARNER

Temperature > 50°F



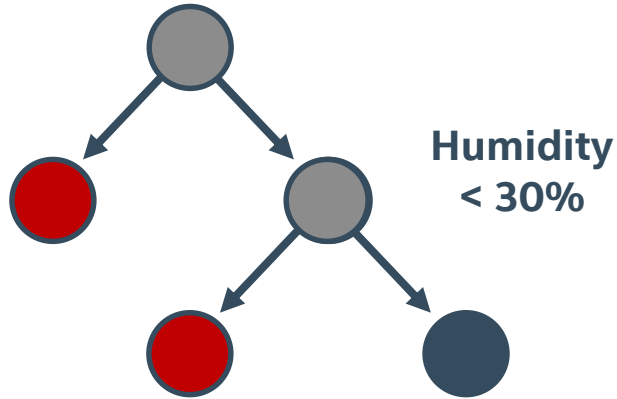
Temperature



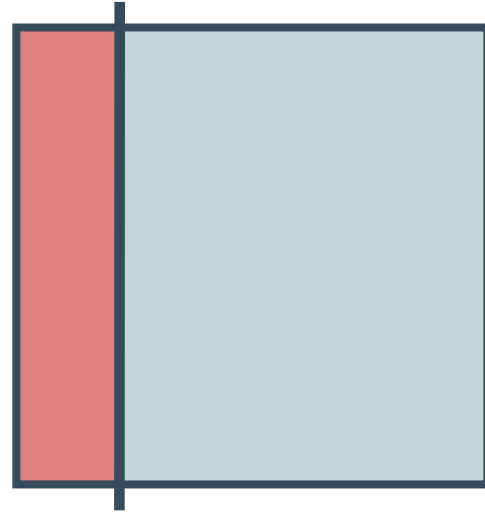


# DECISION STUMP: THE BOOSTING BASE LEARNER

Temperature > 50°F

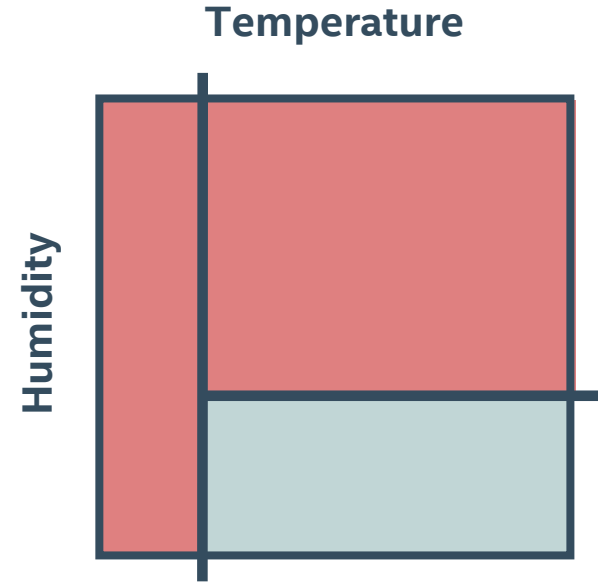
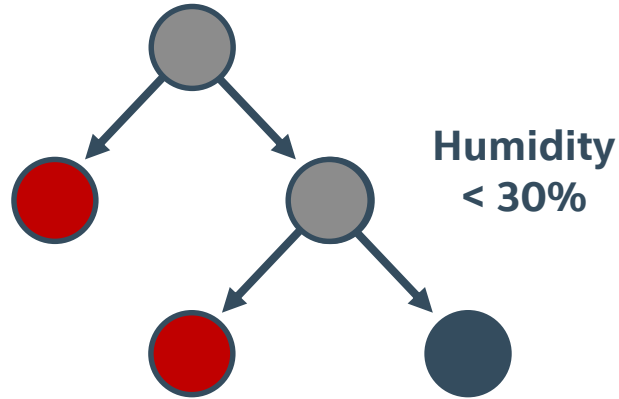


Temperature



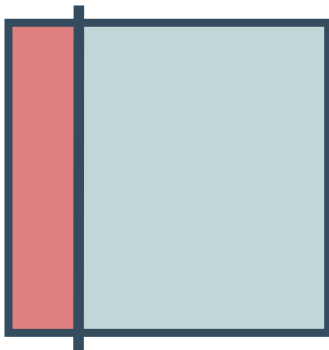
# DECISION STUMP: THE BOOSTING BASE LEARNER

Temperature > 50°F



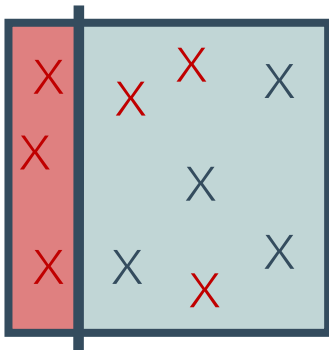
# OVERVIEW OF BOOSTING

Create  
initial  
decision  
stump



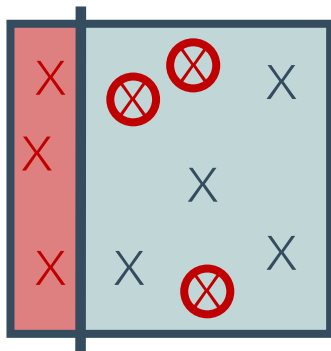
# OVERVIEW OF BOOSTING

Fit to  
data and  
calculate  
residuals



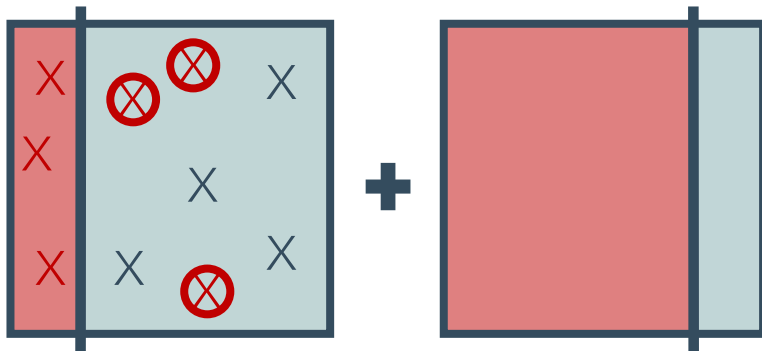
# OVERVIEW OF BOOSTING

Adjust  
weight of  
points



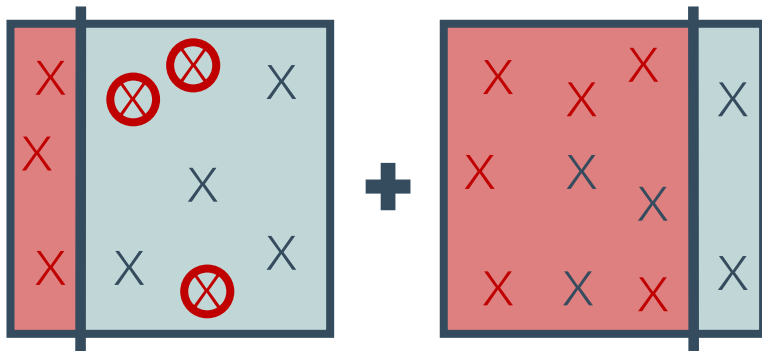
# OVERVIEW OF BOOSTING

Find new  
decision  
stump to fit  
weighted  
residuals



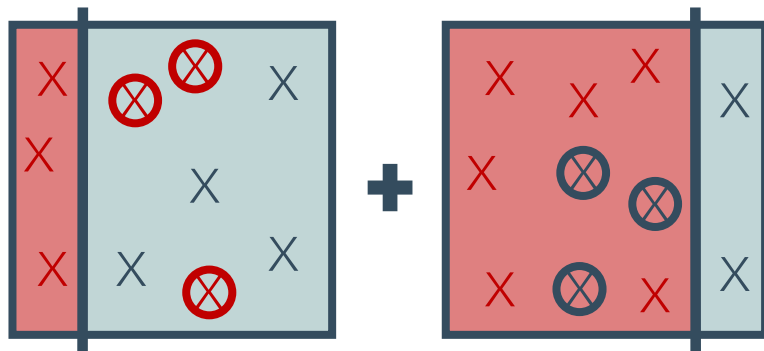
# OVERVIEW OF BOOSTING

Fit new  
decision  
stump to  
current  
residuals



# OVERVIEW OF BOOSTING

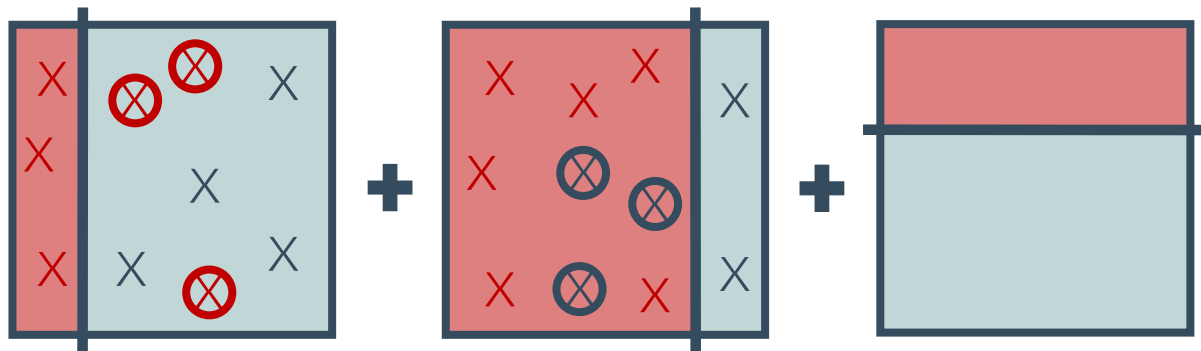
Calculate  
errors and  
weight data  
points





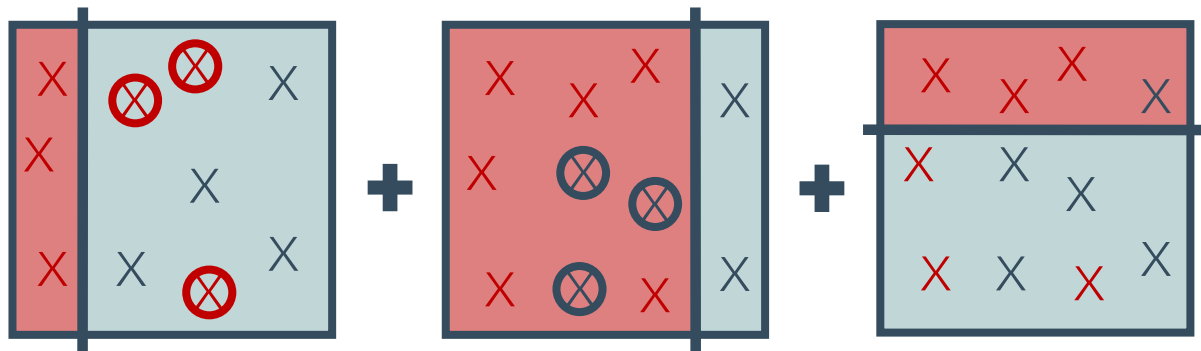
# OVERVIEW OF BOOSTING

Find new  
decision  
stump to fit  
weighted  
residuals

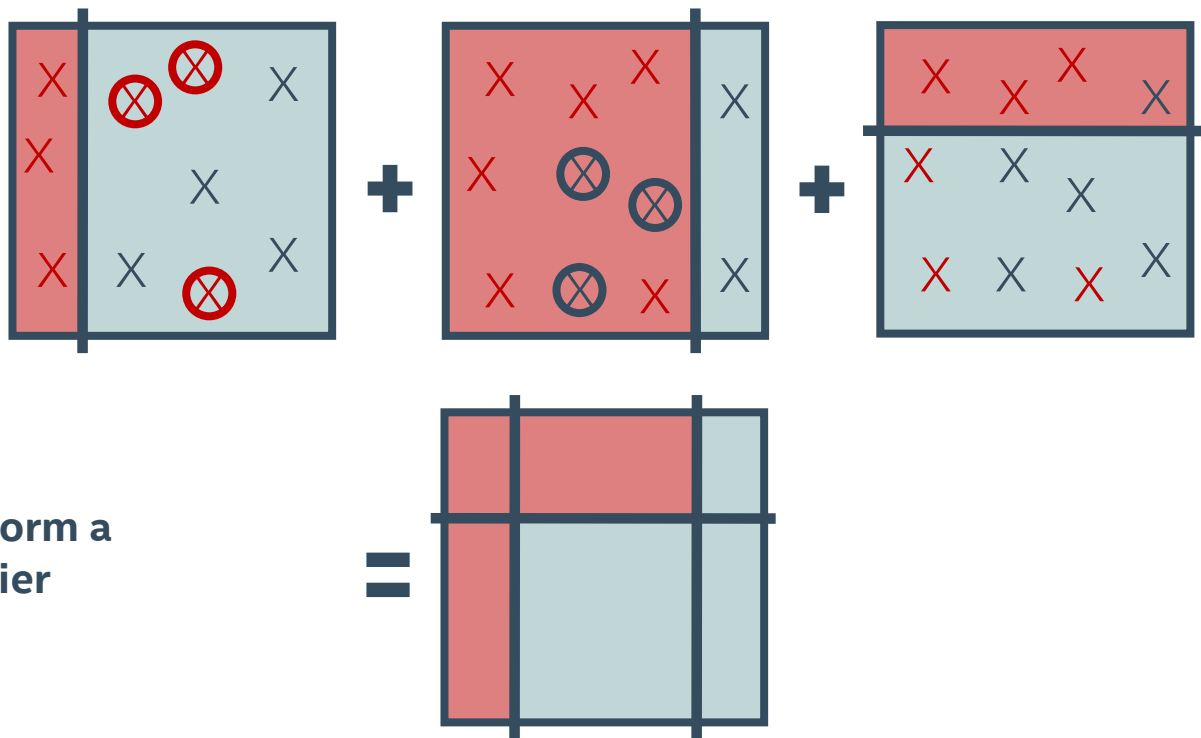


# OVERVIEW OF BOOSTING

Fit new  
decision  
stump to  
current  
residuals

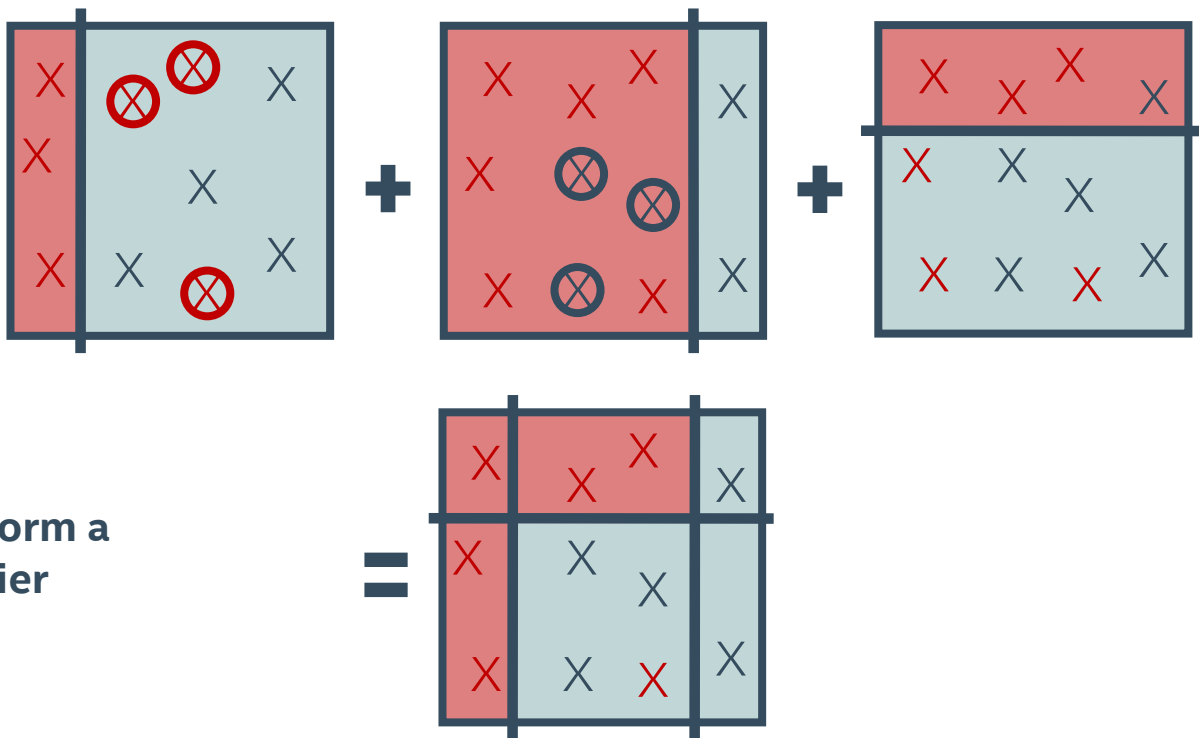


# OVERVIEW OF BOOSTING



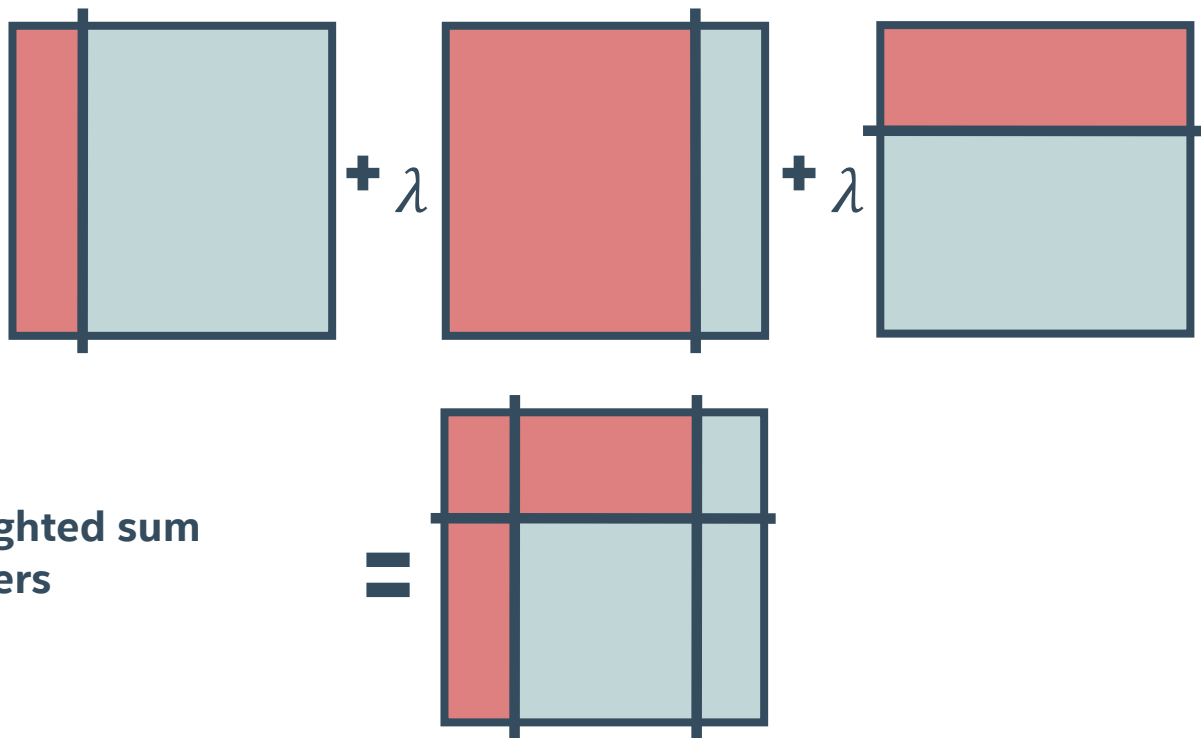
Combine to form a  
single classifier

# OVERVIEW OF BOOSTING



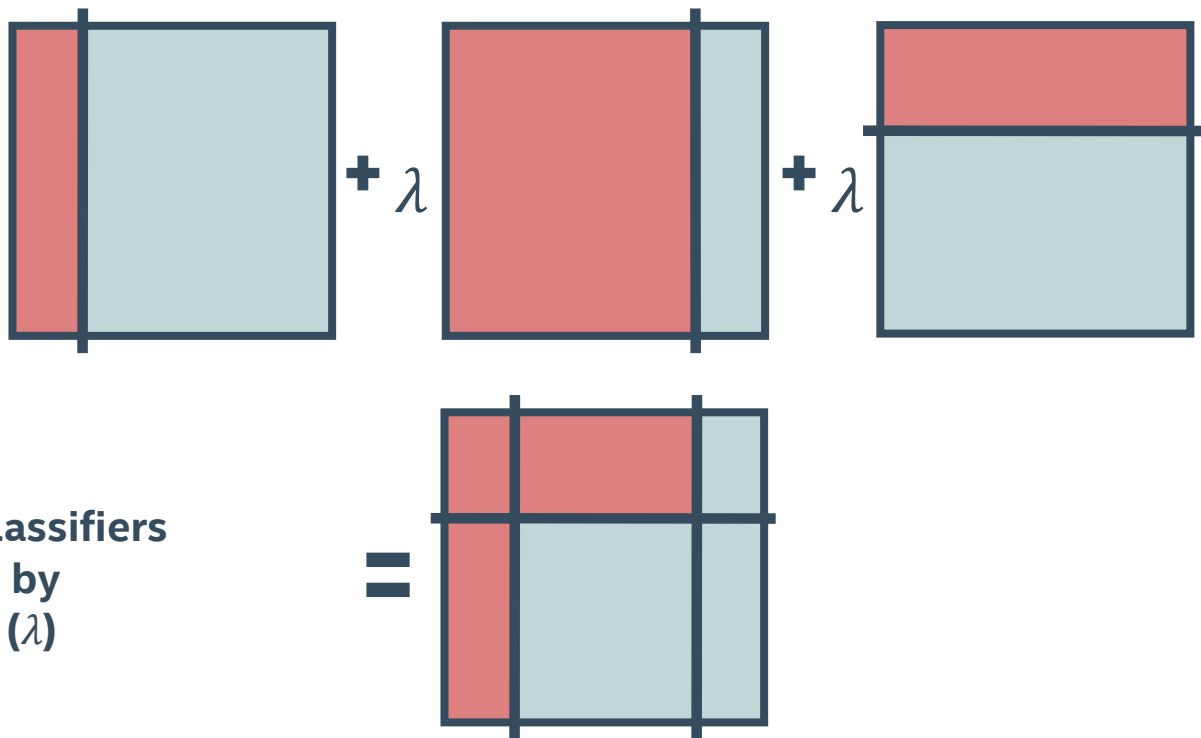
Combine to form a  
single classifier

# OVERVIEW OF BOOSTING

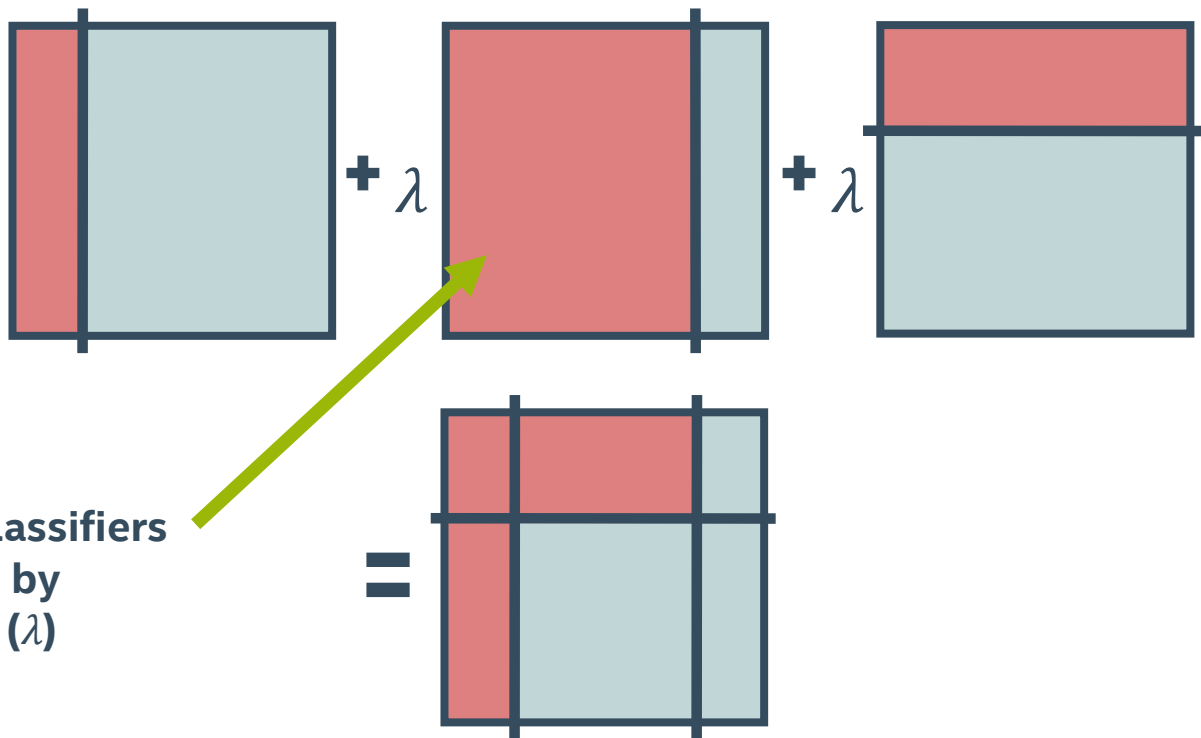


Result is weighted sum  
of all classifiers

# OVERVIEW OF BOOSTING

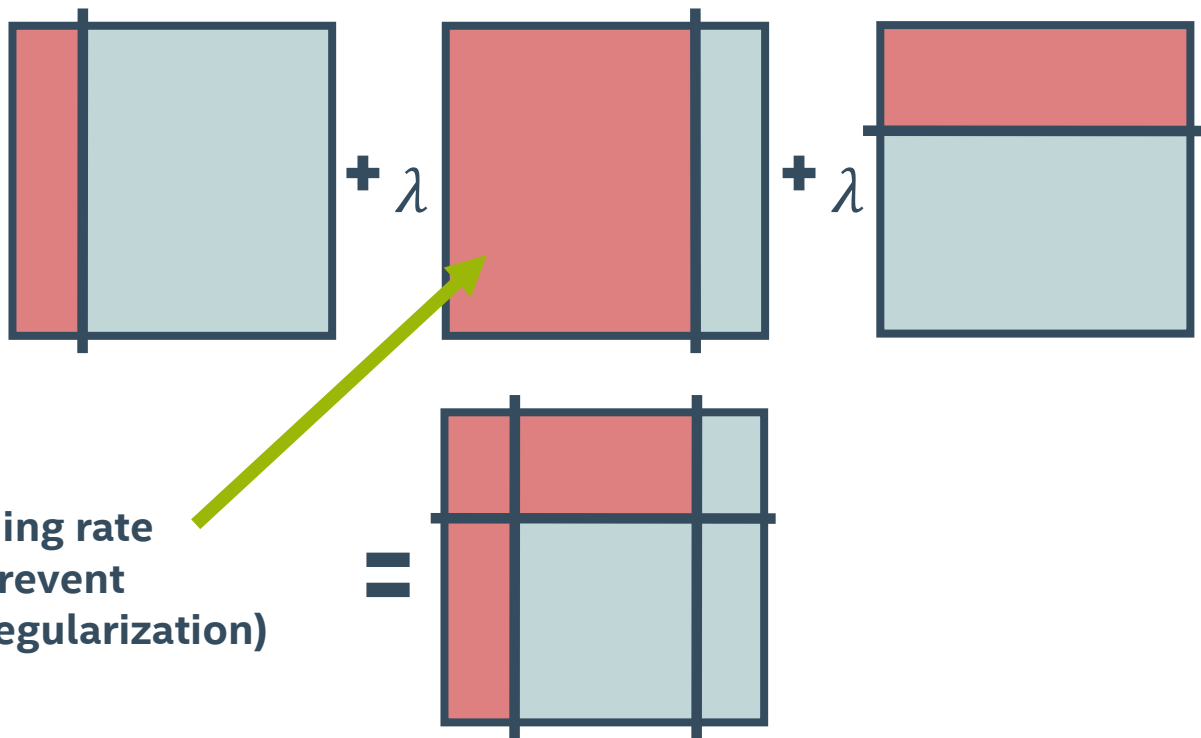


# OVERVIEW OF BOOSTING



Successive classifiers  
are weighted by  
learning rate ( $\lambda$ )

# OVERVIEW OF BOOSTING



Using a learning rate  
< 1.0 helps prevent  
overfitting (regularization)

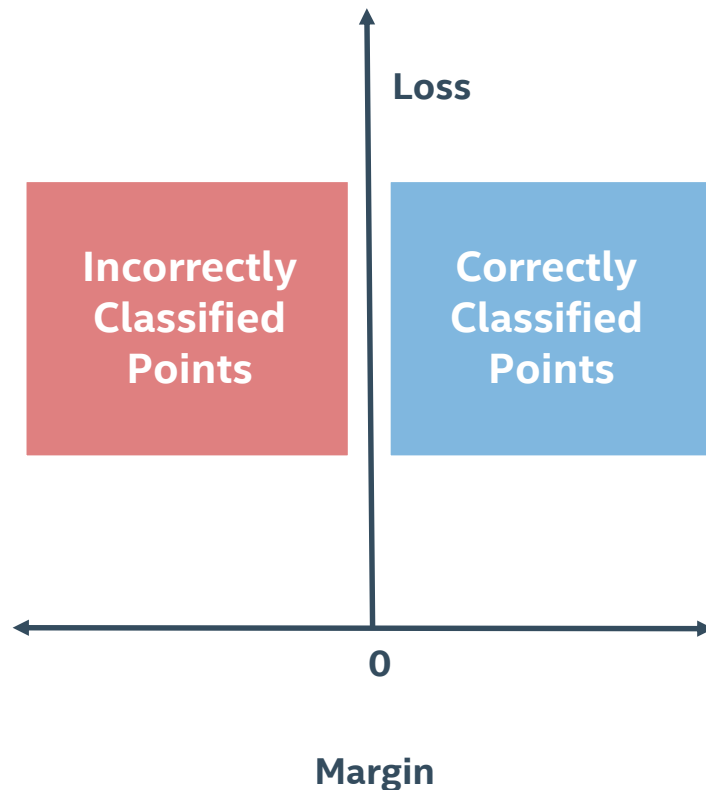


# BOOSTING SPECIFICS

- Boosting utilizes different loss functions
- At each stage, the margin is determined for each point

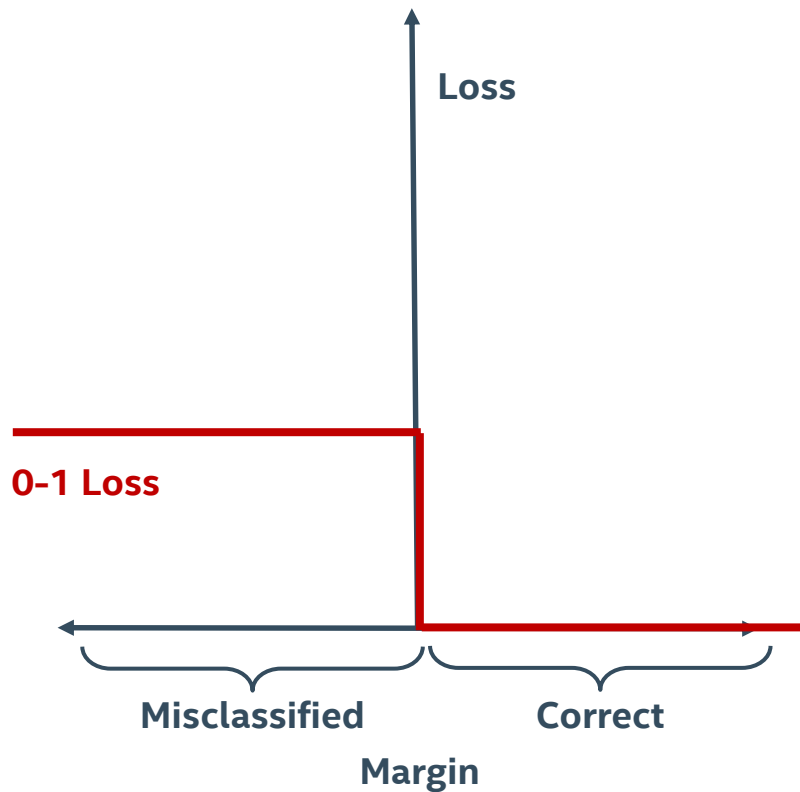
# BOOSTING SPECIFICS

- Boosting utilizes different loss functions
- At each stage, the margin is determined for each point
- Margin is positive for correctly classified points and negative for misclassifications
- Value of loss function is calculated from margin



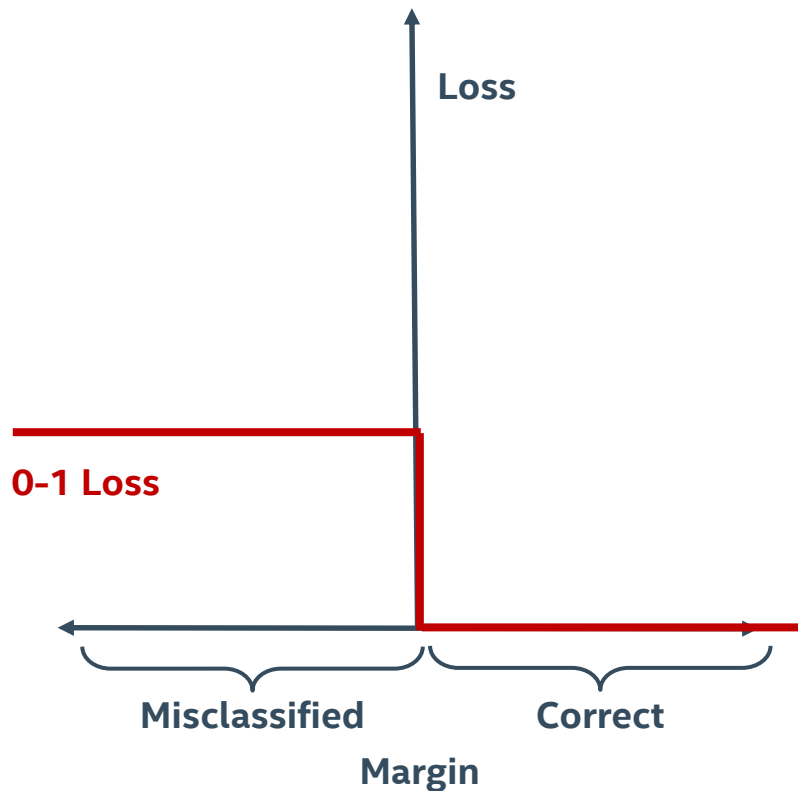
# 0-1 LOSS FUNCTION

- The 0 – 1 Loss multiplies misclassified points by 1



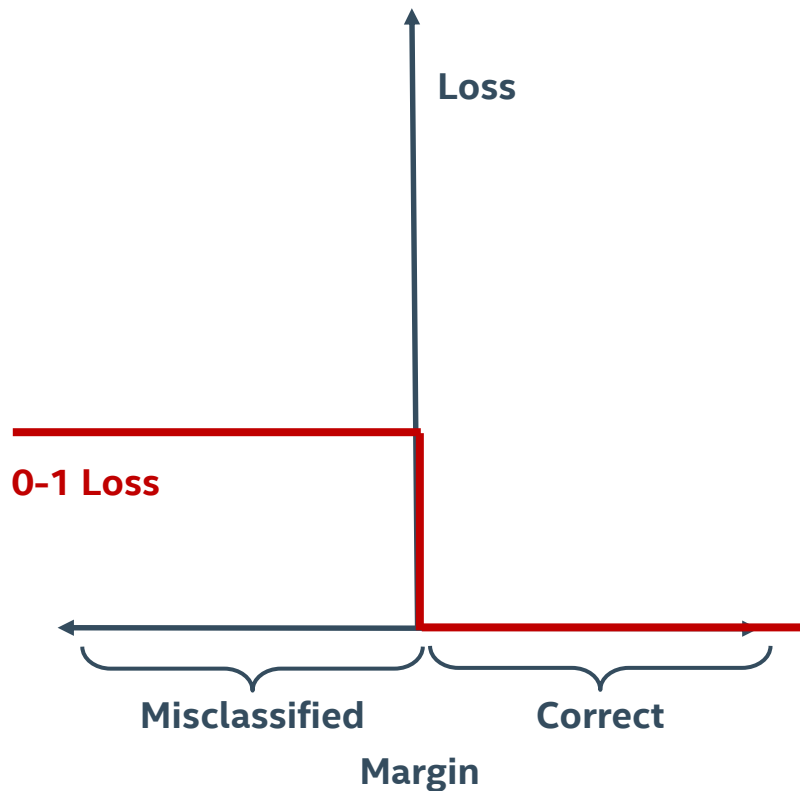
# 0-1 LOSS FUNCTION

- The 0 – 1 Loss multiplies misclassified points by 1
- Correctly classified points are ignored



# 0-1 LOSS FUNCTION

- The 0 – 1 Loss multiplies misclassified points by 1
- Correctly classified points are ignored
- Theoretical "ideal" loss function



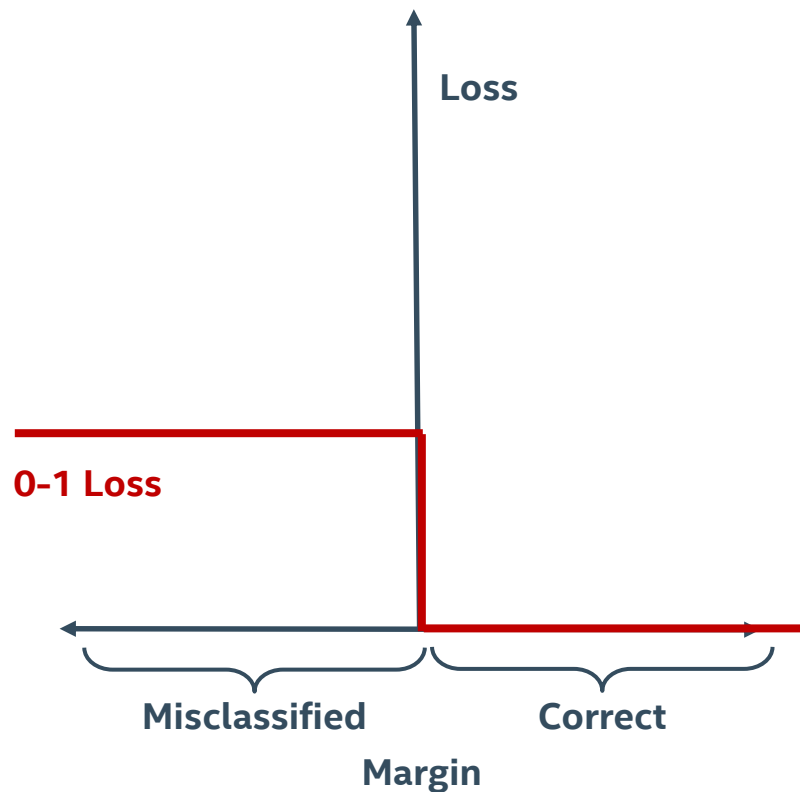
# 0-1 LOSS FUNCTION

- The 0 – 1 Loss multiplies misclassified points by 1
- Correctly classified points are ignored
- Theoretical "ideal" loss function
- Difficult to optimize—non-smooth and non-convex



# ADABOOST LOSS FUNCTION

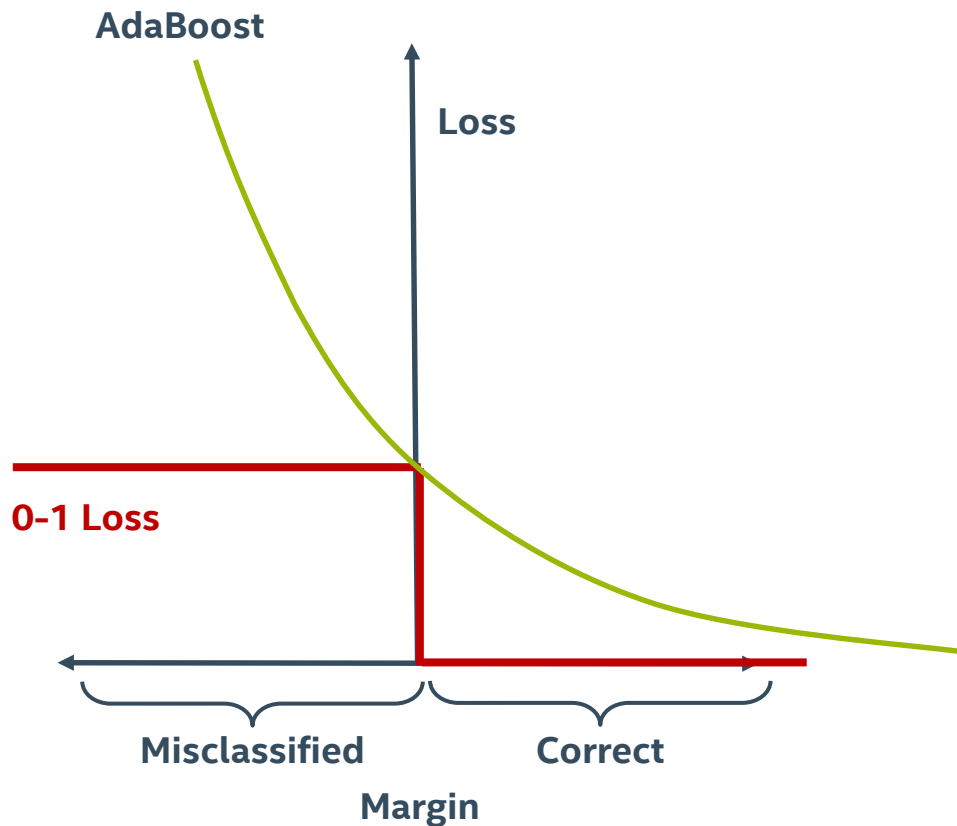
- AdaBoost = Adaptive Boosting



# ADABOOST LOSS FUNCTION

- AdaBoost = Adaptive Boosting
- Loss function is exponential:

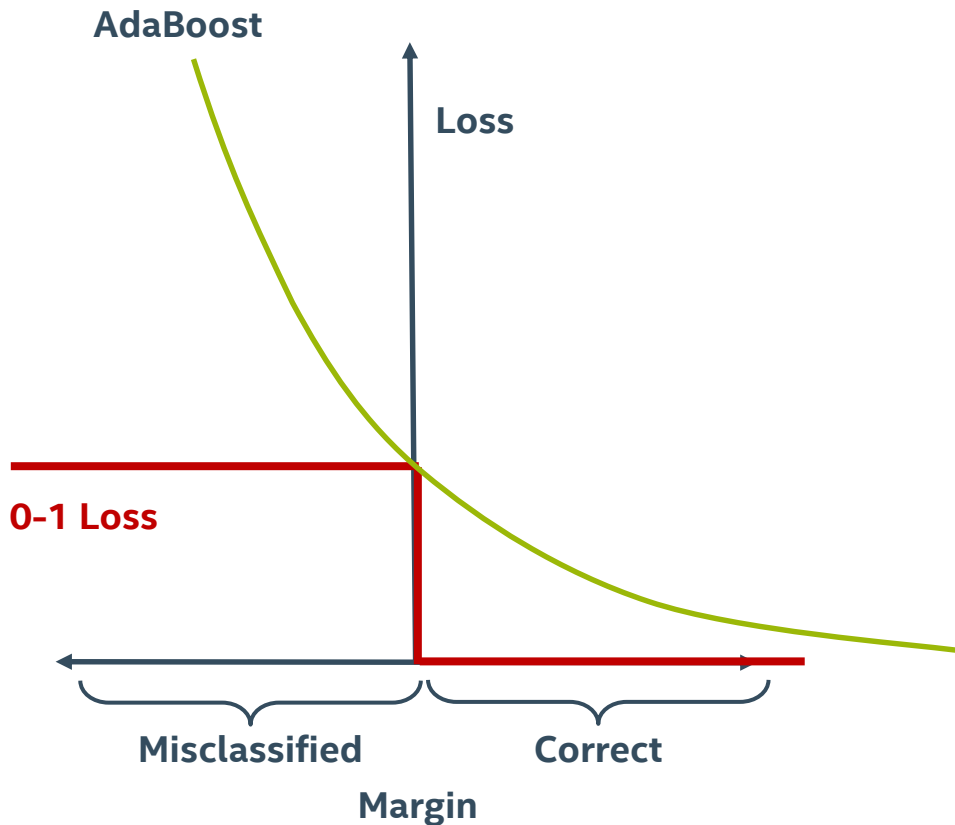
$$e^{-margin}$$





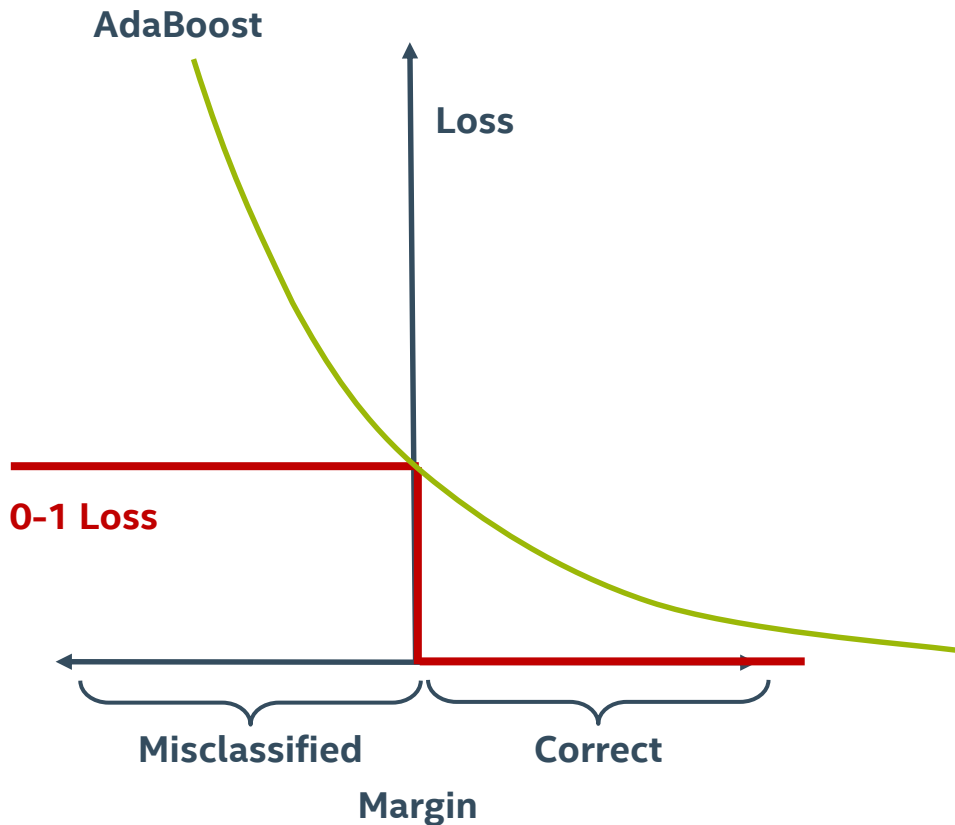
# ADABOOST LOSS FUNCTION

- AdaBoost = Adaptive Boosting
- Loss function is exponential:  
$$e^{-margin}$$
- Makes AdaBoost more sensitive to outliers than other types of boosting



# GRADIENT BOOSTING LOSS FUNCTION

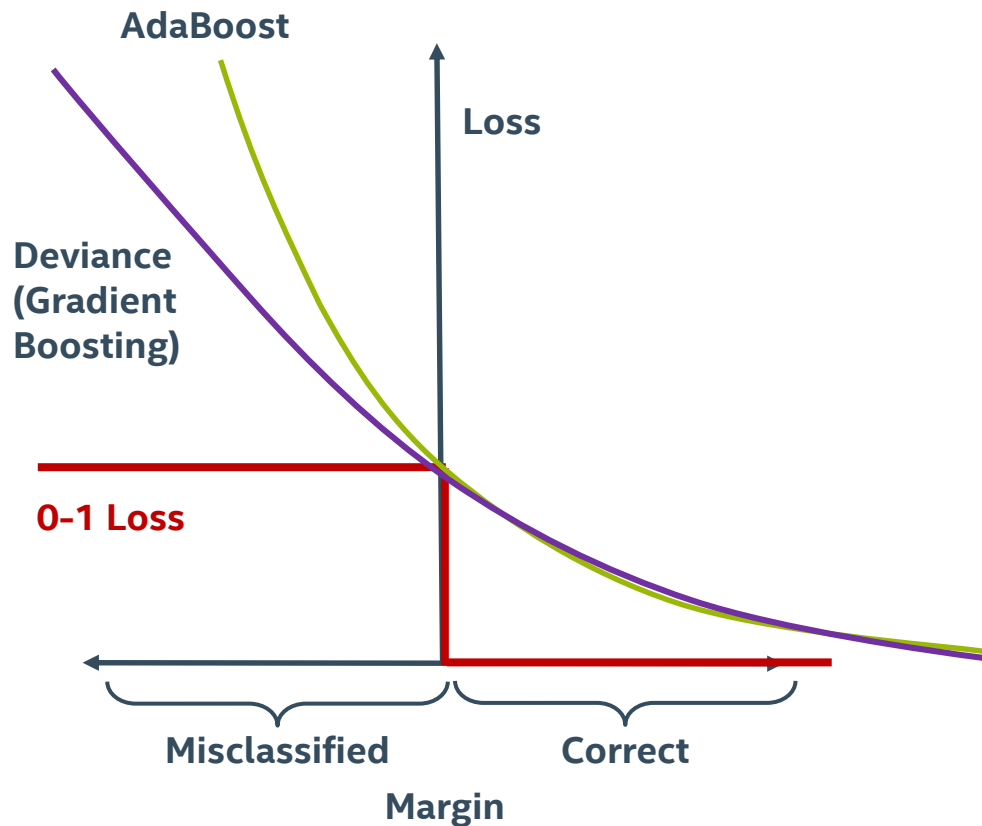
- Generalized boosting method that can use different loss functions
- Common implementation uses binomial log likelihood loss function (deviance):



# GRADIENT BOOSTING LOSS FUNCTION

- Generalized boosting method that can use different loss functions
- Common implementation uses binomial log likelihood loss function (deviance):

$$\log(1 + e^{(-margin)})$$

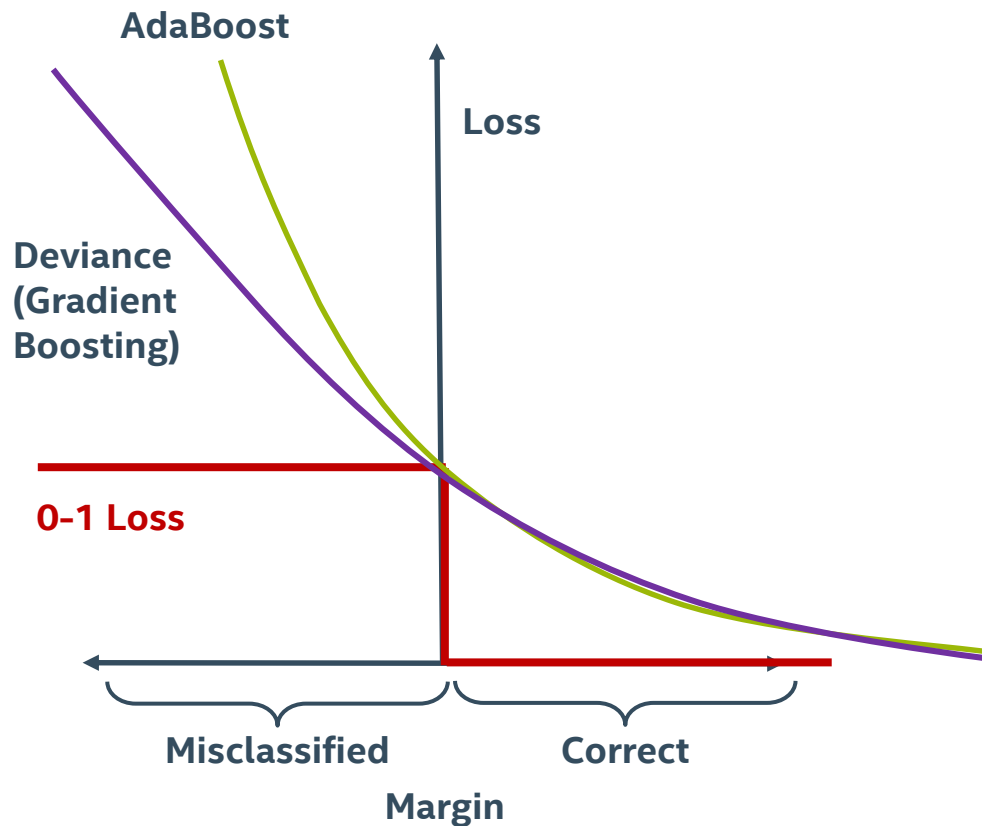


# GRADIENT BOOSTING LOSS FUNCTION

- Generalized boosting method that can use different loss functions
- Common implementation uses binomial log likelihood loss function (deviance):

$$\log(1 + e^{(-margin)})$$

- More robust to outliers than AdaBoost



# BAGGING VS BOOSTING

## BAGGING

- Bootstrapped samples

## BOOSTING

- Fit entire data set

# BAGGING VS BOOSTING

## BAGGING

- Bootstrapped samples
- Base trees created independently

## BOOSTING

- Fit entire data set
- Base trees created successively

# BAGGING VS BOOSTING

## BAGGING

- Bootstrapped samples
- Base trees created independently
- Only data points considered

## BOOSTING

- Fit entire data set
- Base trees created successively
- Use residuals from previous models

# BAGGING VS BOOSTING

## BAGGING

- Bootstrapped samples
- Base trees created independently
- Only data points considered
- No weighting used

## BOOSTING

- Fit entire data set
- Base trees created successively
- Use residuals from previous models
- Up-weight misclassified points



# BAGGING VS BOOSTING

## BAGGING

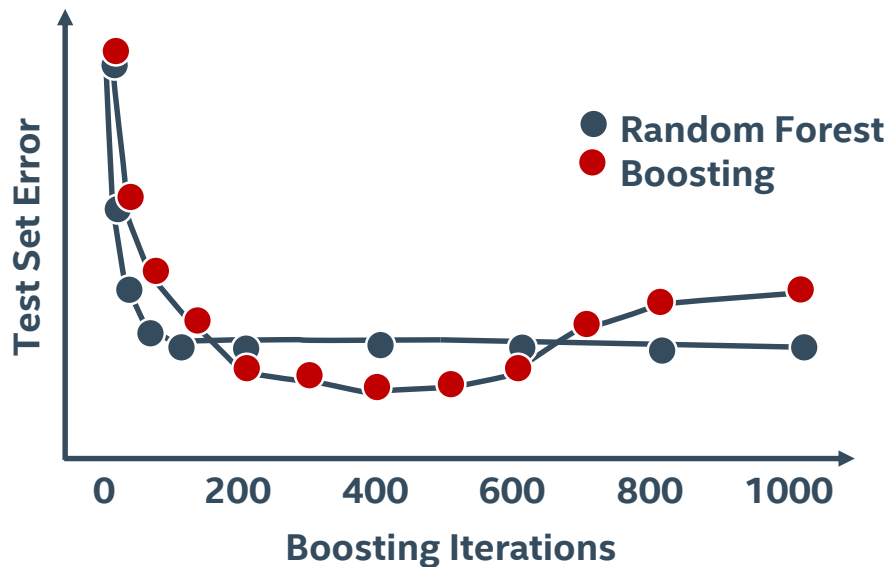
- Bootstrapped samples
- Base trees created independently
- Only data points considered
- No weighting used
- Excess trees will not overfit

## BOOSTING

- Fit entire data set
- Base trees created successively
- Use residuals from previous models
- Up-weight misclassified points
- Beware of overfitting

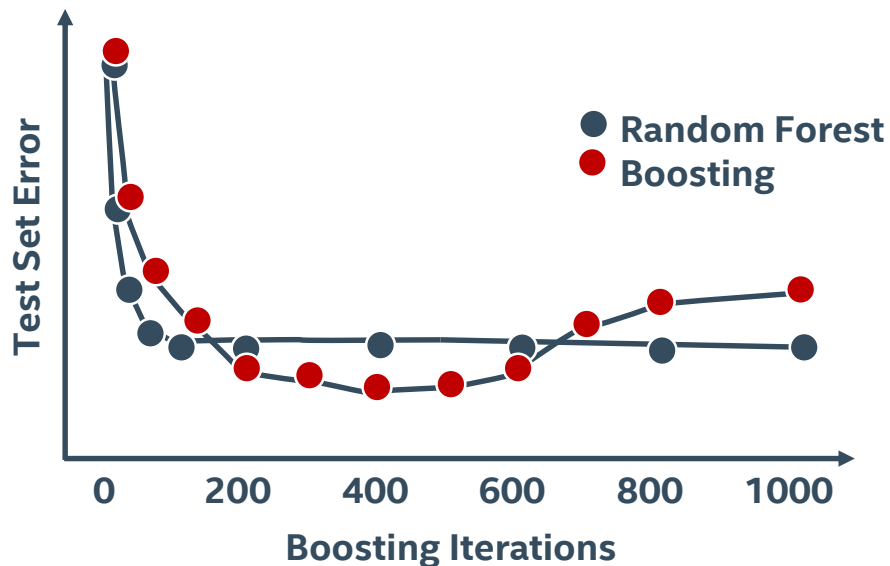
# TUNING A GRADIENT BOOSTED MODEL

- Boosting is additive, so possible to overfit



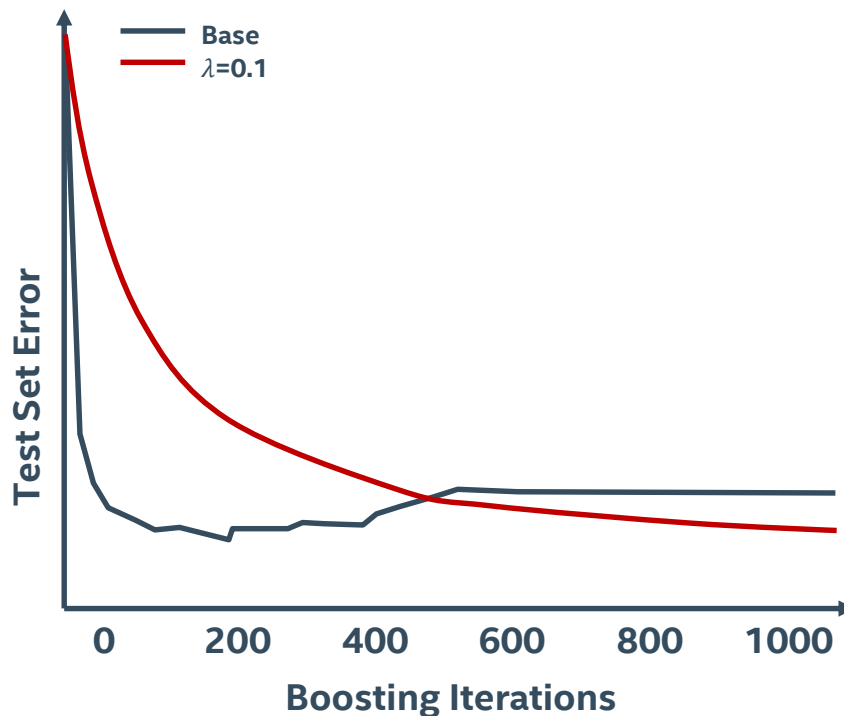
# TUNING A GRADIENT BOOSTED MODEL

- Boosting is additive, so possible to overfit
- Use cross validation to set number of trees



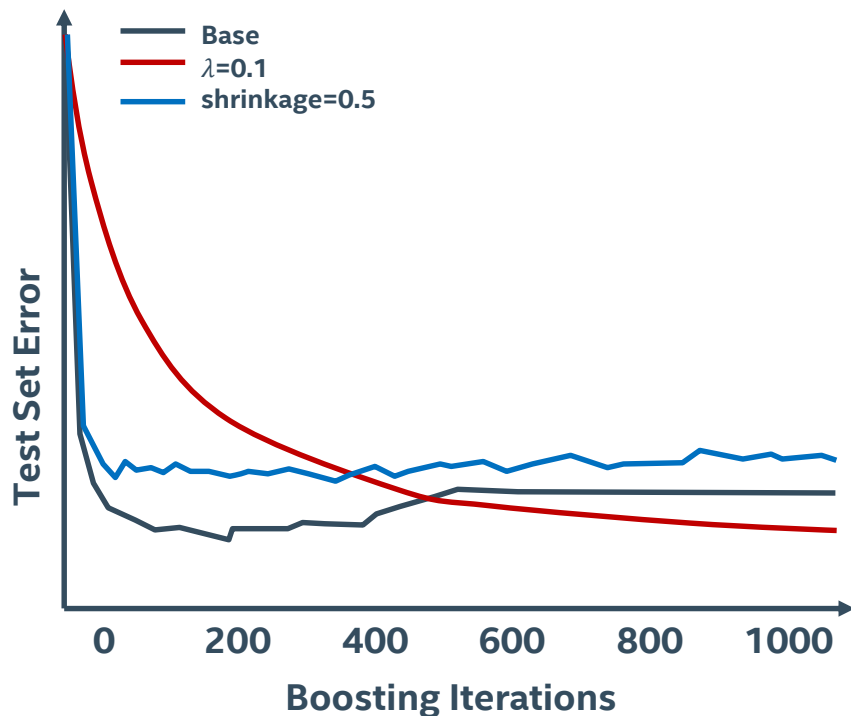
# TUNING A GRADIENT BOOSTED MODEL

- **Learning rate ( $\lambda$ ):** set to  $<1.0$  for regularization. That's also called "shrinkage"



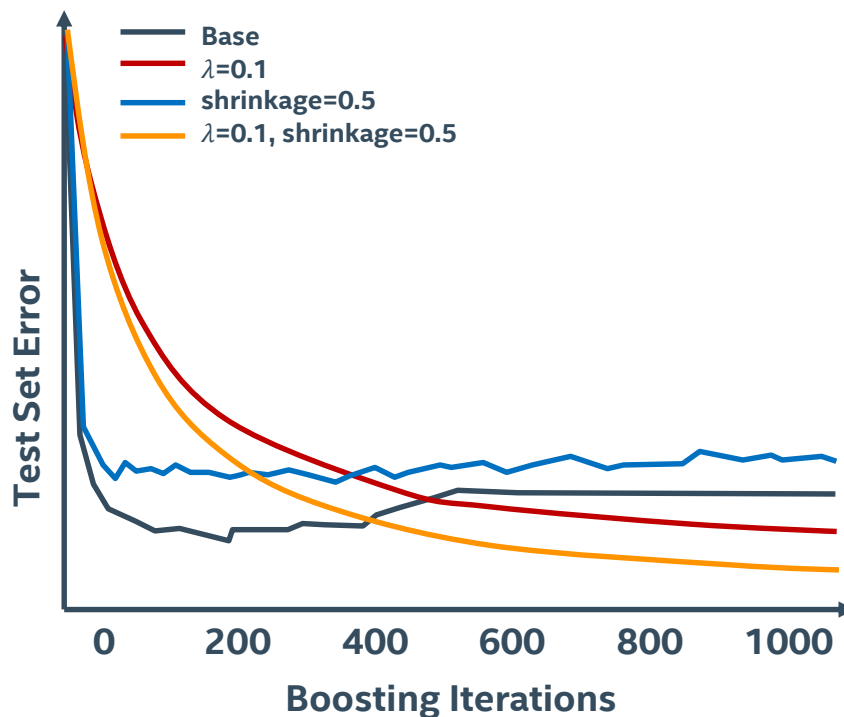
# TUNING A GRADIENT BOOSTED MODEL

- **Learning rate ( $\lambda$ ):** set to  $<1.0$  for regularization. That's also called "shrinkage"
- **Subsample:** set to  $<1.0$  to use fraction of data for base learners (stochastic gradient boosting)



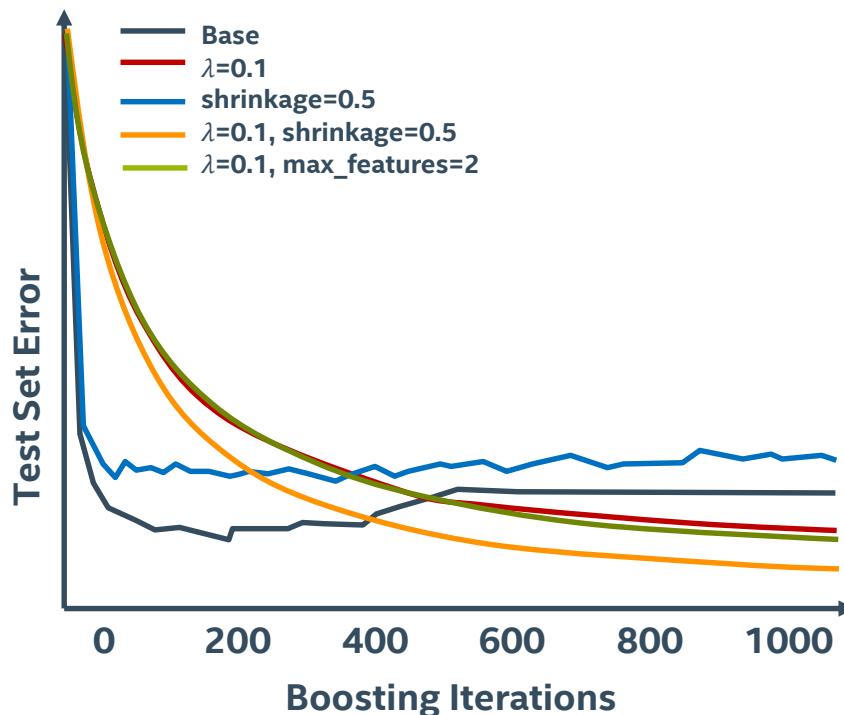
# TUNING A GRADIENT BOOSTED MODEL

- **Learning rate ( $\lambda$ ):** set to  $<1.0$  for regularization. That's also called "shrinkage"
- **Subsample:** set to  $<1.0$  to use fraction of data for base learners (stochastic gradient boosting)



# TUNING A GRADIENT BOOSTED MODEL

- **Learning rate ( $\lambda$ ):** set to  $<1.0$  for regularization. That's also called "shrinkage"
- **Subsample:** set to  $<1.0$  to use fraction of data for base learners (stochastic gradient boosting)
- **Max\_features:** number of features to consider in base learners when splitting.



# GRADIENTBOOSTINGCLASSIFIER: THE SYNTAX

Import the class containing the classification method.

```
from sklearn.ensemble import GradientBoostingClassifier
```



# GRADIENTBOOSTINGCLASSIFIER: THE SYNTAX

Import the class containing the classification method.

```
from sklearn.ensemble import GradientBoostingClassifier
```

Create an instance of the class.

```
GBC = GradientBoostingClassifier(learning_rate=0.1,  
                                max_features=1, subsample=0.5,  
                                n_estimators=200)
```

# GRADIENTBOOSTINGCLASSIFIER: THE SYNTAX

Import the class containing the classification method.

```
from sklearn.ensemble import GradientBoostingClassifier
```

Create an instance of the class.

```
GBC = GradientBoostingClassifier(learning_rate=0.1,  
                                max_features=1, subsample=0.5,  
                                n_estimators=200)
```

Fit the instance on the data and then predict the expected value.

```
GBC = GBC.fit (X_train, y_train)  
y_predict = GBC.predict(X_test)
```

# GRADIENTBOOSTINGCLASSIFIER: THE SYNTAX

Import the class containing the classification method.

```
from sklearn.ensemble import GradientBoostingClassifier
```

Create an instance of the class.

```
GBC = GradientBoostingClassifier(learning_rate=0.1,  
                                max_features=1, subsample=0.5,  
                                n_estimators=200)
```

Fit the instance on the data and then predict the expected value.

```
GBC = GBC.fit (X_train, y_train)  
y_predict = GBC.predict(X_test)
```

Tune with cross-validation. Use `GradientBoostingRegressor` for regression.

# ADABOOSTCLASSIFIER: THE SYNTAX

Import the class containing the classification method.

```
from sklearn.ensemble import AdaBoostClassifier  
from sklearn.tree import DecisionTreeClassifier
```

# ADABOOSTCLASSIFIER: THE SYNTAX

Import the class containing the classification method.

```
from sklearn.ensemble import AdaBoostClassifier  
from sklearn.tree import DecisionTreeClassifier
```

Create an instance of the class.

```
ABC = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(),  
                          learning_rate=0.1, n_estimators=200)
```

# ADABOOSTCLASSIFIER: THE SYNTAX

Import the class containing the classification method.

```
from sklearn.ensemble import AdaBoostClassifier  
from sklearn.tree import DecisionTreeClassifier
```

Create an instance of the class.

```
ABC = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(),  
                          learning_rate=0.1, n_estimators=200)
```



base learner  
can be set  
manually

# ADABOOSTCLASSIFIER: THE SYNTAX

Import the class containing the classification method.

```
from sklearn.ensemble import AdaBoostClassifier  
from sklearn.tree import DecisionTreeClassifier
```

Create an instance of the class.

```
ABC = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(),  
                          learning_rate=0.1, n_estimators=200)
```

can also set  
max depth here



# ADABOOSTCLASSIFIER: THE SYNTAX

Import the class containing the classification method.

```
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
```

Create an instance of the class.

```
ABC = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(),
                          learning_rate=0.1, n_estimators=200)
```

Fit the instance on the data and then predict the expected value.

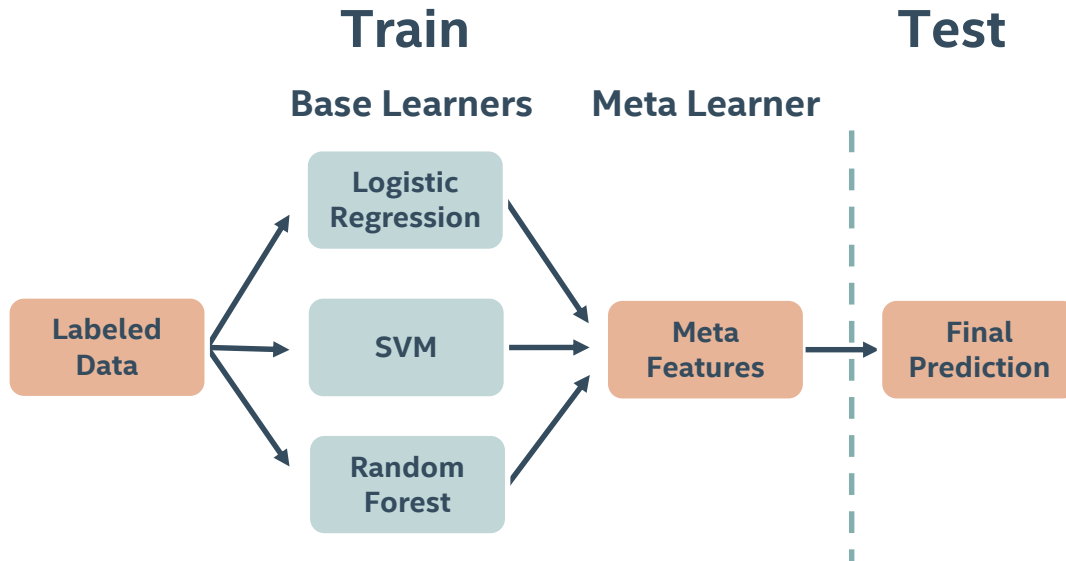
```
ABC = ABC.fit(X_train, y_train)
y_predict = ABC.predict(X_test)
```

Tune with cross-validation. Use `AdaBoostRegressor` for regression.



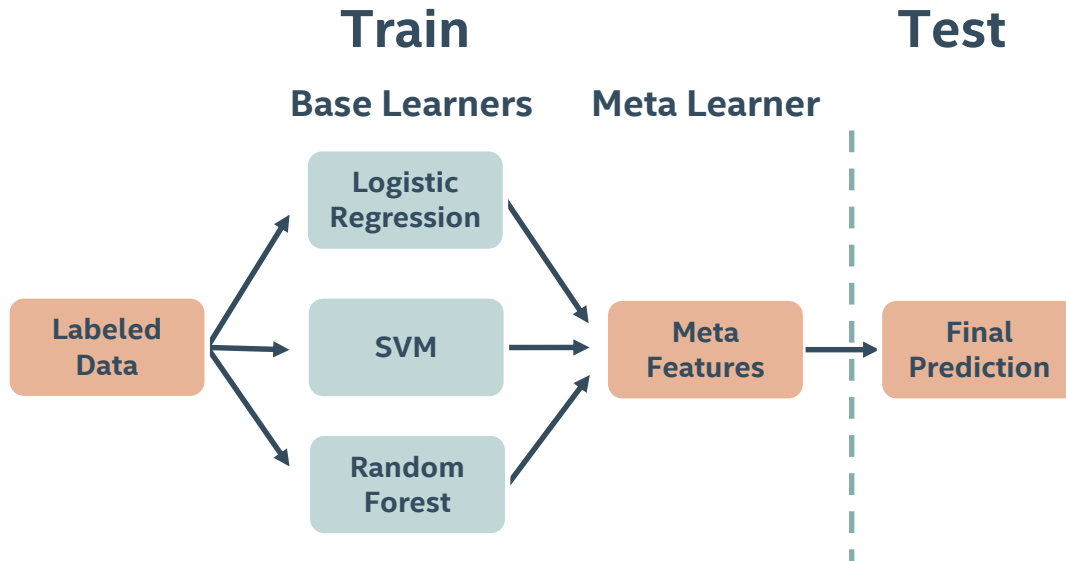
# STACKING: COMBINING HETEROGENEOUS CLASSIFIERS

- Models of any kind combined to create stacked model



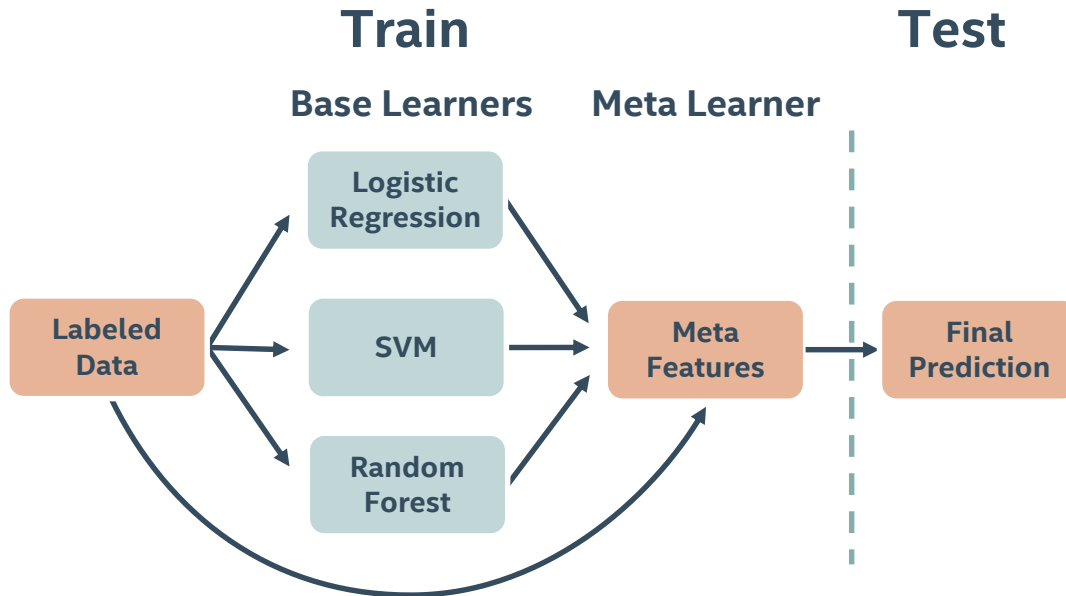
# STACKING: COMBINING HETEROGENEOUS CLASSIFIERS

- Models of any kind combined to create stacked model
- Like bagging but not limited to decision trees



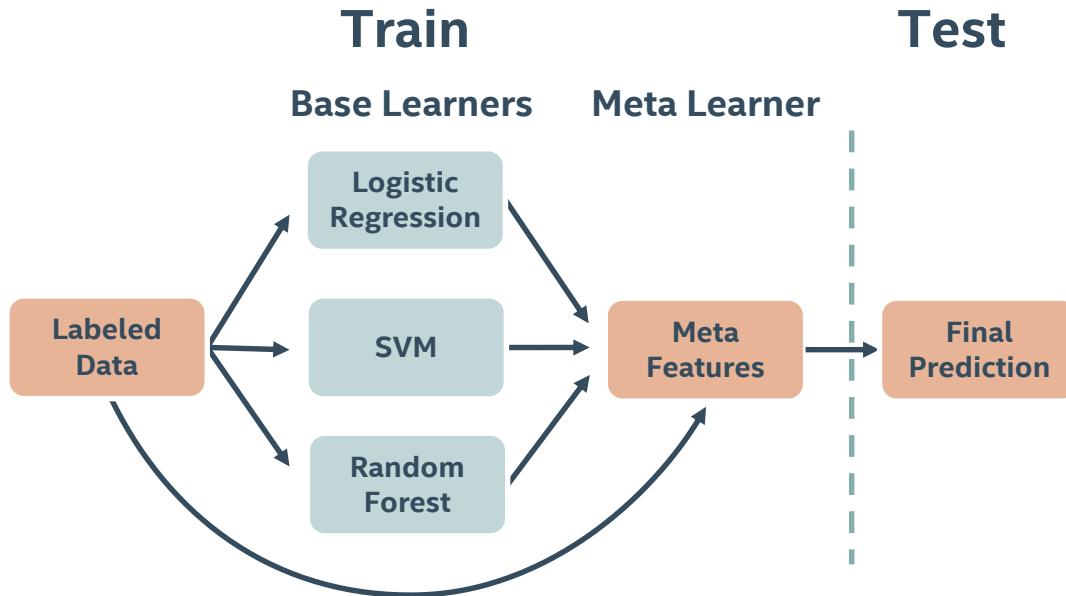
# STACKING: COMBINING HETEROGENEOUS CLASSIFIERS

- Models of any kind combined to create stacked model
- Like bagging but not limited to decision trees
- Output of base learners creates features, can recombine with data



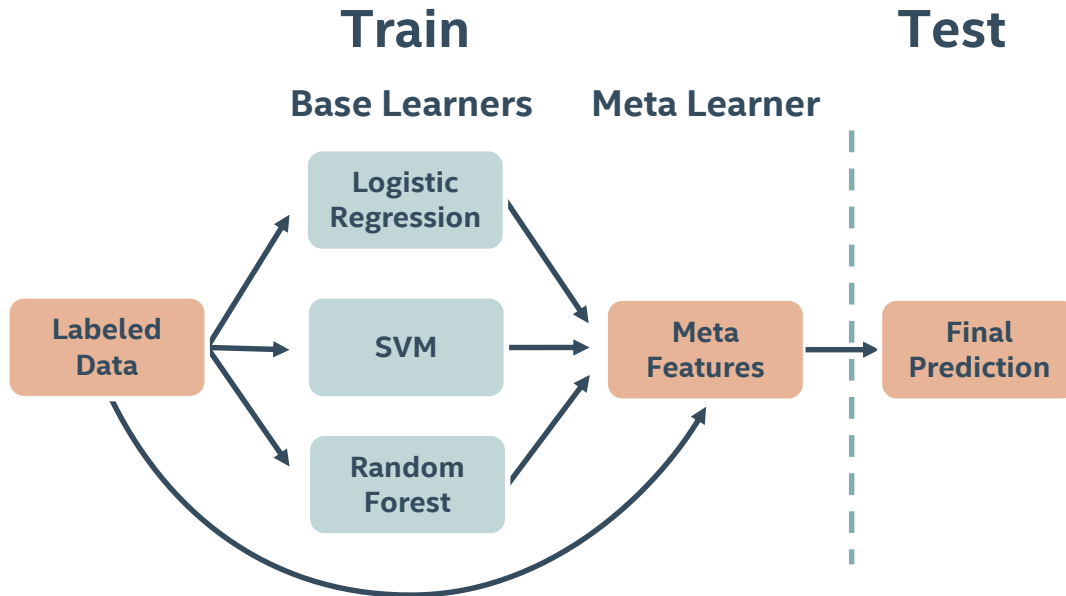
# STACKING: COMBINING HETEROGENEOUS CLASSIFIERS

- Output of base learners can be combined via majority vote or weighted



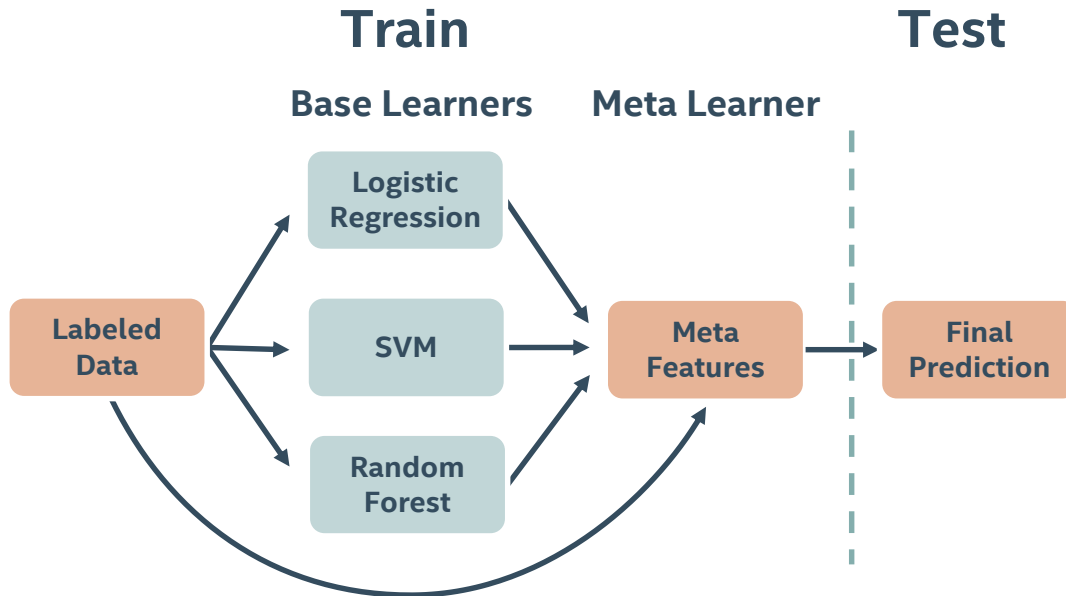
# STACKING: COMBINING HETEROGENEOUS CLASSIFIERS

- Output of base learners can be combined via majority vote or weighted
- Additional hold-out data needed if meta learner parameters are used



# STACKING: COMBINING HETEROGENEOUS CLASSIFIERS

- Output of base learners can be combined via majority vote or weighted
- Additional hold-out data needed if meta learner parameters are used
- Be aware of increasing model complexity



# VOTINGCLASSIFIER: THE SYNTAX

Import the class containing the classification method.

```
from sklearn.ensemble import VotingClassifier
```

# VOTINGCLASSIFIER: THE SYNTAX

Import the class containing the classification method.

```
from sklearn.ensemble import VotingClassifier
```

Create an instance of the class.

```
VC = VotingClassifier(estimator_list, voting='hard',  
                      weights=estimator_weight_list)
```



# VOTINGCLASSIFIER: THE SYNTAX

Import the class containing the classification method.

```
from sklearn.ensemble import VotingClassifier
```

Create an instance of the class.

```
VC = VotingClassifier(estimator_list, voting='hard',  
                      weights=estimator_weight_list)
```



**list of fitted  
models and  
how to combine**

# VOTINGCLASSIFIER: THE SYNTAX

Import the class containing the classification method.

```
from sklearn.ensemble import VotingClassifier
```

Create an instance of the class.

```
VC = VotingClassifier(estimator_list, voting='hard',  
                      weights=estimator_weight_list)
```

Fit the instance on the data and then predict the expected value.

```
VC = V.fit(X_train, y_train)  
y_predict = VC.predict(X_test)
```

Tune with an ADDITIONAL LEVEL of cross-validation or hold-out set.

